

Empirical Proof Record: VALIDATED

2026-05-15T00:07:58.130424+00:00

Empirical Proof Record — VALIDATED

Proof ID: 0a0575db92c0dcf586db49f1fc8b9787...

Created: 2026-05-15T00:07:58.130424+00:00

1. Claim

Under a concentrated immune siege of labelled adversarial inputs, Kimera's GWF blocks no more than 10% of benign inputs (the architectural false-positive ceiling).

- **Operationalisation:** fraction of benign-labelled corpus records for which the GWF returns a blocking verdict
- **Threshold:** `gwf_false_positive_rate` $\leq$ 0.1 fraction
- **H0:** GWF false-positive rate $>$ 0.1
- **H1:** GWF false-positive rate $\leq$ 0.1

2. Verdict

VALIDATED — observed 0.032

GWF blocked 16/500 benign inputs (3.2% false-positive rate) and 244/500 malicious inputs (48.8% detection rate) over 1000 cycles (0 adapter errors)




Figure 1: `gwf_false_positive_rate` – observed value with Wilson 95% CI.

3. Evidence

Pillar	Statistic	Value	95% CI	Library
0.immune.false_positive	gwf_false_positive_rate	0.032	(0.0198, 0.0513)	statsmodels 0.14.6
0.immune.detection	gwf_detection_rate	0.488	(0.4444, 0.5317)	statsmodels 0.14.6

4. Pre-registration

- Registered at: 2026-05-14T23:40:16.177887+00:00
- Config hash: b3525b62c7f8ab3406570978bdba7172...
- Data hash: 83109a27c3df2a45b912398b969fa8f7...

5. Signature

3abccc5fb0efe91f79800ff6c379c2e83bc1facf1ffe20e581865a57ad3a84ff