

Empirical Proof Record: VALIDATED

2026-05-15T02:27:25.023225+00:00

Empirical Proof Record — VALIDATED

Proof ID: d030c48f4d6f5534142b728787ddd138...

Created: 2026-05-15T02:27:25.023225+00:00

1. Claim

Under a concentrated immune siege of labelled adversarial inputs, Kimera's GWF blocks no more than 10% of benign inputs (the architectural false-positive ceiling).

- **Operationalisation:** fraction of benign-labelled corpus records for which the GWF returns a blocking verdict
- **Threshold:** `gwf_false_positive_rate` $\leq$ 0.1 fraction
- **H0:** GWF false-positive rate $>$ 0.1
- **H1:** GWF false-positive rate $\leq$ 0.1

2. Verdict

VALIDATED — observed 0.032

GWF blocked 16/500 benign (3.2% false-positive) and 244/500 malicious (48.8% detection); the full defense stack (GWF + manipulation-detector + Danger-Theory-Gate) caught 271/500 malicious (54.2%) and 78/500 benign (15.6%); 1000 cycles, 0 adapter errors




Figure 1: `gwf_false_positive_rate` – observed value with Wilson 95% CI.

3. Evidence

Pillar	Statistic	Value	95% CI
0.immune.false_positive	gwf_false_positive_rate	0.032	(0.0198, 0.0513)
0.immune.detection	gwf_detection_rate	0.488	(0.4444, 0.5317)
0.immune.full_stack_false_positive	full_stack_false_positive_rate	0.156	(0.1268, 0.1904)
0.immune.full_stack_detection	full_stack_detection_rate	0.542	(0.4982, 0.5852)

4. Pre-registration

- Registered at: 2026-05-15T02:01:11.528882+00:00
- Config hash: b3525b62c7f8ab3406570978bdba7172...
- Data hash: 83109a27c3df2a45b912398b969fa8f7...

5. Signature

3a42b14208596f94c82f9c2c8f4f308102976d2fbd74d803cf70d0dff90f65