

Empirical Proof Record: VALIDATED

2026-05-23T23:11:02.667582+00:00

Empirical Proof Record — VALIDATED

Proof ID: 22ec9910f4466322ad81b5df9d360711...

Created: 2026-05-23T23:11:02.667582+00:00

1. Claim

Under a concentrated immune siege of labelled adversarial inputs, Kimera's GWF blocks no more than 10% of benign inputs (the architectural false-positive ceiling).

- **Operationalisation:** fraction of benign-labelled corpus records for which the GWF returns a blocking verdict
- **Threshold:** `gwf_false_positive_rate` $\leq$ 0.1 fraction
- **H0:** GWF false-positive rate $>$ 0.1
- **H1:** GWF false-positive rate $\leq$ 0.1

2. Verdict

VALIDATED — observed 0.0

GWF blocked 0/20 benign (0.0% false-positive) and 8/20 malicious (40.0% detection); the gwf direct target has no other defense layers — full-stack rates equal the GWF rates; 40 cycles, 0 adapter errors




Figure 1: `gwf_false_positive_rate` – observed value with Wilson 95% CI.

3. Evidence

Pillar	Statistic	Value	95% CI
0.immune.false_positive	gwf_false_positive_rate	0.0	(0.0000, 0.1611)
0.immune.detection	gwf_detection_rate	0.4	(0.2188, 0.6134)
0.immune.full_stack_false_positive	full_stack_false_positive_rate	0.0	(0.0000, 0.1611)
0.immune.full_stack_detection	full_stack_detection_rate	0.4	(0.2188, 0.6134)

4. Pre-registration

- Registered at: 2026-05-23T23:09:39.429937+00:00
- Config hash: 2dadbb261d8fdef7ef0810b289205c3a...
- Data hash: 83109a27c3df2a45b912398b969fa8f7...

5. Signature

08e305a2163fce545c6c4ab256afd2778fe66f58d4c9753abe4de34024d5087d